

VeritasGraph Studio


Build Governed AI Agents. On-Prem.

Most enterprise AI demos answer "what does the doc say?" — VeritasGraph Studio answers the question that actually blocks production: "who is violating this policy — and can you prove it, on our own infrastructure?"

OPEN SOURCE

AIR-GAPPED

AUDIT-READY



One Workspace. Ten Capabilities.

VeritasGraph Studio is a single workspace where enterprise teams **build, test, evaluate, and govern AI agents** — all running locally on your own models. No data leaves your VPC. Here's how a platform team would walk through it.



Agents

Compose role-focused agents with capability toggles and per-agent context budgets.



Tools

Register and validate callable APIs, MCP servers, and retrievers before any agent uses them.



Knowledge

Manage indexed corpora with hybrid search and graph-linking for high attribution confidence.




Guardrails

Enforce hard-block and review policies — PII redaction, toxicity monitoring — on every agent turn.



Agents




Compose role-focused agents — orchestrator, reviewer, assistant — and wire exactly which capabilities each one gets. Every agent ships with a **per-agent context budget** and an orchestration map that updates in real time as you configure it.

→ Knowledge Graph

→ Tools + Memory

→ Guardrails + Data Logging

Tools & Knowledge



Tools

Register callable tools — internal APIs, MCP servers, retrievers — with schema and auth. Run **smoke tests** before any agent can use them, so nothing ships un-validated. The registry shows connected vs. disabled tools at a glance.

Knowledge

Manage indexed corpora (Policy Corpus, Product Wiki) with **hybrid search plus graph-linking**. Tune chunk size and overlap to maximize attribution confidence for every retrieval.

AGENT

Based on the provided knowledge graph and documents, I cannot determine their sex.

I can only confirm their **names** (Alice and Bob), their **roles** (Alice: Accounts Payable and Vendor Management; Bob: Accounts Payable), and their **actions** (Alice: approved Acme Corp contract; Bob: reviewed budget).

YOU

My email is test_user@domain.com and my credit card number is 4111 1111 1111 1111. Analyze this text.

AGENT

recalled 6 turn(s)

Knowledge Graph

20 nodes · 40 edges · seeds: Alice, Rule R7, Bob

Veritasroom budget

567/600 tokens · kept 4 · dropped 7

Tools

available: Graph Rtriever, Web Search, Design Beard, VeritasGraph MCP
SearXNG Search, n8n Workflows, Knowledge Graph · Invoked 1 (ok 1, failed 0)

Guardrails (out)

0 redaction(s)

CAPABILITY 04

Guardrails


Define **hard-block and review policies** — PII redaction, toxicity monitoring — enforced on every agent turn. A live incident count provides real-time audit visibility across all active agents.

PII Filter	Toxicity Monitor	Live Incident Count
Enforced on every agent turn	Review or block policy triggers	Audit-ready counter

CAPABILITY 05


CAPABILITY 06

Memory & Data



Memory


Give agents **session (short-term)** and **long-term memory scopes**.
Built-in hygiene checks flag stale, duplicated, or conflicting memories before they corrupt agent reasoning.



Data

Connect operational data — SQL, files, events, object storage — and gate it with **pre-index quality checks** to reduce noisy retrieval and improve answer quality.


Evaluation



Benchmark every agent for **relevance, faithfulness, latency, and policy compliance**. Track a rolling quality trend across runs — so you catch regressions before release, not after.

- Evaluation is a first-class capability, not an afterthought. Every agent ships with a quality baseline.

Fine-Tune & Playground




Fine-Tune

Queue domain-adaptation jobs on curated data slices. **Safety gates are enforced on every fine-tuned checkpoint** – no model update reaches production without passing compliance checks.

Playground

Test any agent live against its local model. Watch the full orchestration pipeline – **guardrails in, memory, knowledge graph, context budget, tools, guardrails out, data log** – with citations on every answer.

Knowledge Graph



Build a VeritasGraph knowledge graph from your documents and reason over it with **verifiable citations and an explicit reasoning path**. Ingest a document, explore nodes and edges, and ask multi-hop questions — every answer is grounded and auditable.

Ingest

Documents → indexed graph nodes

Explore

Visualize nodes, edges, and relationships

Reason

Multi-hop queries with PROV-O lineage

Governed. Auditable. Air-Gapped.

Guardrails, evaluation, and data-quality gates are **first-class**. Every answer carries citations and PROV-O lineage for audit. The whole thing runs air-gapped on your infrastructure — build, test, and govern agents in one place instead of stitching five tools together.



Governed

Policies enforced at every agent turn



Auditable

Citations + PROV-O lineage on every answer



Air-Gapped

Runs entirely on your infrastructure



★ Open source and free. Star the repo, spin up the Studio, and tell us which agent you'd build first.

